

NLP501 - SESSION 03

# Vector Space Models

*Words as Vectors in Semantic Space*

# Learning Objectives

1. Understand word vectors in semantic space

4. Calculate cosine similarity between vectors

2. Build co-occurrence matrices from text

5. Apply PCA for dimensionality reduction

3. Apply TF-IDF weighting schemes

6. Implement document search systems

**Exercise 1 Released Today!** Sentiment Analysis and Vector Spaces (Due: Session 5)

# Session Outline

**Part 1:** Quick Review - Session 02 Naïve Bayes

**Part 5:** Cosine Similarity

**Part 2:** Word Representations

**Part 6:** Dimensionality Reduction (PCA)

**Part 3:** Co-occurrence Matrices

**Part 7:** Applications

**Part 4:** TF-IDF Weighting

**Part 8:** Summary and Lab Practice

PART 1

# Quick Review

*Session 02 - Naïve Bayes Classifier*

# Session 02 Review: Naïve Bayes

## Key Formulas

$$P(c|d) = P(d|c) \times P(c) / P(d)$$

$$P(w_1, w_2, \dots | c) = \prod P(w_i | c)$$

$$P(w | c) = (\text{count}(w, c) + \alpha) / (N^c + \alpha |V|)$$

## Key Concepts

**Bayes Theorem:** Update prior beliefs with evidence

**Naïve Assumption:** Features are conditionally independent

**Laplacian Smoothing:** Handle unseen words (add  $\alpha$ )

**Log-space:** Prevent underflow with log probabilities

**Key Insight:** Naïve Bayes is simple, fast, and works well for text classification despite the "naïve" assumption!

# From Counts to Vectors: The Journey

## Session 01: Logistic Regression

Feature extraction, Sigmoid function, Gradient descent optimization

## Session 02: Naïve Bayes

Word frequencies, Conditional probabilities, Bayesian classification

## Session 03: Vector Space Models

Words as vectors, Semantic similarity, Document search

### Key Question for Today:

*"How do we represent the MEANING of words?"*

PART 2

# Word Representations

*From Symbols to Vectors*

"You shall know a word by the company it keeps" — J.R. Firth (1957)

# One-Hot Encoding

## Definition

Vector with dimension =  $|V|$ , all zeros except one position = 1

**Example:**  $V = [\text{cat}, \text{dog}, \text{happy}, \text{sad}]$

$\text{cat} = [1, 0, 0, 0]$

$\text{dog} = [0, 1, 0, 0]$

$\text{happy} = [0, 0, 1, 0]$

$\text{sad} = [0, 0, 0, 1]$

## Problems

**No Similarity:** All word pairs are orthogonal

**High Dimensional:** Vector size = vocabulary size

**Sparse:** Only one non-zero element

**No Semantics:**  $\cos(\text{cat}, \text{dog}) = 0$

**Key Issue:** "cat" and "dog" should be similar (both animals), but one-hot gives  $\cos(v_{\text{cat}}, v_{\text{dog}}) = 0$

# The Distributional Hypothesis

*"Words that occur in similar contexts tend to have similar meanings"*

— Zellig Harris (1954), J.R. Firth (1957)

## Example: Similar Contexts

"The CAT sat on the mat"  
"The DOG sat on the mat"  
"A cute CAT is sleeping"  
"A cute DOG is sleeping"

✓ "cat" and "dog" share contexts → similar vectors

## Example: Different Contexts

"I feel HAPPY today"  
"The ALGORITHM converged"  
"She is HAPPY with results"  
"The ALGORITHM is efficient"

X "happy" and "algorithm" → different vectors

## 💡 Key Insight: We can learn word meaning from context!

Instead of defining meanings manually, we let data tell us which words are similar based on usage patterns.

# Two Types of Vector Space Models

## Sparse Vectors

Count-based Methods

- Long vectors (10,000+ dims)
- Most values = 0
- Based on co-occurrence counts

Methods: Term-Doc Matrix, TF-IDF

This Session's Focus

## Dense Vectors

Prediction-based Methods

- Short vectors (50-300 dims)
- All values non-zero
- Learned from prediction tasks

Methods: Word2Vec, GloVe, FastText

Session 06: Word Embeddings

Both approaches capture semantic similarity, but dense vectors are more efficient for downstream tasks

PART 3

# Co-occurrence Matrices

---

*Building Word Vectors from  
Context*

# Term-Document Matrix

## Definition

A matrix where:

- **Rows = Words (terms)**
- **Columns = Documents**
- **Cell [i,j] = Count of word i in doc j**

## Term-Document Matrix

Term	D1	D2	D3	D4
i	1	0	0	1
love	1	0	0	1
machine	1	1	2	0
learning	1	1	2	1
deep	0	0	1	1
is	0	1	0	0
awesome	0	1	0	0
and	0	0	1	0

## Example Corpus

- D1: "I love machine learning"  
D2: "Machine learning is awesome"  
D3: "Deep learning and machine learning"  
D4: "I love deep learning"

$$v_{\text{machine}} = [1, 1, 2, 0] \cdot v_{\text{learning}} = [1, 1, 2, 1]$$

**Key:** Each word becomes a vector over documents. Similar words appear in similar documents!

# Word-Word Co-occurrence Matrix

## Definition

A square matrix where:

- **Rows & Columns = Words**
- **Cell  $[i,j]$  = Count of word j near word i**
- **Symmetric:  $M[i,j] = M[j,i]$**

## Context Window

Defines "nearness" - typically k words on each side

**Window k=2:** 2 words left + 2 words right

## Word-Word Matrix (k=1)

Sentence: I love deep learning

Pair: (I, love), (love, I), (love, deep), (deep, love), (deep, learning), (learning, deep)

	I	love	deep	learning
I	0	1	0	0
love	1	0	1	0
deep	0	1	0	1
learning	0	0	1	0

**Benefit:** Captures semantic relationships directly between words

# Context Window: Example

Sentence: "The quick brown fox jumps over the lazy dog"



Target: "fox" · Context window k=2

## Context Words for "fox" (k=2)

Left: quick, brown

Right: jumps, over

Context = {quick, brown, jumps, over}

## Window Size Effects

- **Small (k=1-2):** Syntactic relationships
- **Medium (k=4-5):** Balanced
- **Large (k=10+):** Topical/semantic

## Key Insight

Window size is a hyperparameter. Small windows capture grammar, large windows capture topics.

PART 4

# TF-IDF Weighting

---

*Measuring Term Importance*

# The Problem with Raw Counts

## Document About Machine Learning

"THE basics of MACHINE LEARNING are THE foundation for THE field... THE algorithms..."

**15**

"the" count

**3**

"ML" count

## The Problem

- Common words like "the" have high counts
- Important topic words have low counts
- Raw counts don't reflect importance!

**Issue:** "the" dominates but tells us nothing about the document's topic

## Solution: TF-IDF

Weight terms by how important they are to a document relative to the entire corpus. Downweight common words, upweight distinctive words.

# Term Frequency (TF)

## Definition

Measures how frequently a term appears in a document.

$$tf(t, d) = f(t,d)$$

Count of term t in document d

## Common Variants

### Raw Count

$$tf(t, d) = f(t,d)$$

### Log Normalization

$$tf(t, d) = 1 + \log(f(t,d))$$

### Double Normalization

$$tf(t, d) = 0.5 + 0.5 \times f(t,d) / \max(fd)$$

## Example Calculation

Document: "machine learning is great. machine learning is powerful. machine learning! "

	machine	learning	is
Raw TF	3	3	2
Boolean TF	1	1	1
Log TF	$1 + \log(3) \approx 2.10$	$1 + \log(3) \approx 2.10$	$1 + \log(2) \approx 1.69$
Augmented TF	$0.5 + 0.5 \times (3/3) = 1.0$	$0.5 + 0.5 \times (3/3) = 1.0$	$0.5 + 0.5 \times (2/3) \approx 0.83$

Note: Log dampens the effect of high frequency.  $TF=3 \rightarrow 2.10$ , not  $3 \times$ .

# Inverse Document Frequency (IDF)

## Definition

Measures how **rare** or **common** a term is across all documents.

$$\text{idf}(t) = \log(N / \text{df}(t))$$

N = Total docs    df = Docs with t

## Example: N = 1,000,000 docs

word	df(t)	Calculation	idf(t)
the	10,000	$\log(10000/10000)$	0.00
computer	1,000	$\log(10000/1000)$	1.00
algorithm	100	$\log(10000/100)$	2.00
transformer	10	$\log(10000/10)$	3.00

Note: IDF = 0 when term appears in ALL documents

## Intuition

- Rare terms → High IDF (informative)
- Common terms → Low IDF (less useful)

# TF-IDF: The Complete Formula

$$\text{tf-idf}(t, d) = \text{tf}(t, d) \times \text{idf}(t)$$

$$\text{tf}(t, d) = 1 + \log(f_{t,d})$$

$$\text{idf}(t) = \log(N / df_t)$$

## What TF-IDF Captures

- High TF + High IDF: Important, distinctive
- High TF + Low IDF: Common (downweighted)
- Low TF + High IDF: Rare but not frequent
- Low TF + Low IDF: Not important

## Example: "neural" in ML doc

$$\text{tf} = 1 + \log(5) = 2.61$$

$$\text{idf} = \log(1000/50) = 1.30$$

$$\text{tf-idf} = 2.61 \times 1.30 = 3.39$$

## 💡 Key Insight

TF-IDF balances local importance (TF) with global rarity (IDF). High scores = terms that are frequent in document AND rare across corpus.

# TF-IDF: Python Implementation

```
import numpy as np
from collections import Counter
import math

def compute_tf(doc):
    """Tính Term Frequency cho một document"""
    word_counts = Counter(doc)
    max_count = max(word_counts.values())
    tf = {word: count / max_count for word, count in word_counts.items()}
    return tf

def compute_idf(corpus):
    """Tính IDF cho toàn bộ corpus"""
    N = len(corpus)
    # Đếm số documents chứa mỗi từ
    df = Counter()
    for doc in corpus:
        unique_words = set(doc)
        for word in unique_words:
            df[word] += 1
    # Tính IDF
    idf = {word: math.log(N / count) for word, count in df.items()}
    return idf

def compute_tfidf(doc, idf):
    """Tính TF-IDF vector cho một document"""
    tf = compute_tf(doc)
    tfidf = {word: tf[word] * idf.get(word, 0) for word in tf}
    return tfidf
```

## Using sklearn (Recommended)

```
from sklearn.feature_extraction.text import TfidfVectorizer
# Khởi tạo TfidfVectorizer
vectorizer = TfidfVectorizer(
    lowercase=True,           # Chuyển về lowercase
    stop_words='english',    # Loại bỏ stop words
    max_df=0.9,              # Bỏ từ xuất hiện >90% docs
    min_df=1,                # Giữ từ xuất hiện ít nhất 1 doc
    ngram_range=(1, 2)        # Unigrams và bigrams
)

# Fit và transform
tfidf_matrix = vectorizer.fit_transform(corpus)
```

## sklearn Parameters

- **max\_df:** Ignore terms in > X% docs
- **min\_df:** Ignore terms in < X docs
- **ngram\_range:** (1,2) for unigrams+bigrams

**Note:** sklearn uses slightly different formulas with L2 normalization by default.

PART 5

# Cosine Similarity

---

*Measuring Semantic Distance*

# Why Not Euclidean Distance?

## The Problem with Euclidean

Two documents about same topic:

**Doc A (short):** "machine learning" → [1, 1, 0, 0]

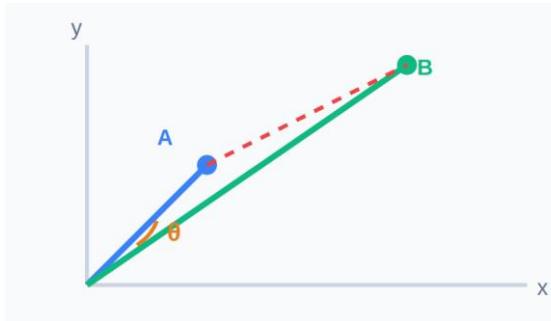
**Doc B (long):** "machine learning ×3" → [3, 3, 0, 0]

### ✗ Euclidean Distance

$$d(A, B) = \sqrt{[(3-1)^2 + (3-1)^2]} = \sqrt{8} \approx 2.83$$

**Problem:** Same topic but large distance!

## Geometric Interpretation



**Euclidean**

Line length

**Cosine**

Angle  $\theta$

✓ Same direction = Same topic, regardless of length!

# Cosine Similarity: The Formula

$$\cos(A, B) = (A \cdot B) / (\|A\| \times \|B\|)$$

$A \cdot B = \sum a_i b_i$  (Dot product)

$\|A\| = \sqrt{(\sum a_i^2)}$  (L2 norm)

## Value Range



- **$\cos = 1$** : Identical direction
- **$\cos = 0$** : Orthogonal (unrelated)
- **$\cos = -1$** : Opposite direction

## Key Properties

- **Length-invariant:** Only direction matters
- **Symmetric:**  $\cos(A, B) = \cos(B, A)$
- **Efficient:**  $O(d)$  computation
- **Non-negative for TF-IDF:** Range  $[0, 1]$

## Cosine Distance

$$d(A, B) = 1 - \cos(A, B)$$

Converts similarity to distance. Range  $[0, 2]$ . Used in clustering algorithms.

# Cosine Similarity: Example

## Document A

"I love machine learning"

$$\mathbf{A} = [1, 1, 1, 1, 0, 0]$$

## Document B

"I love deep learning"

$$\mathbf{B} = [1, 1, 0, 1, 1, 0]$$

## Document C

"NLP is fun"

$$\mathbf{C} = [0, 0, 0, 0, 0, 1]$$

Vocabulary: (I, love, machine, learning, deep, NLP)

## Step-by-Step Calculation

**cos(A, B):**

$$A \cdot B = 1 \times 1 + 1 \times 1 + 1 \times 0 + 1 \times 1 + 0 \times 1 + 0 \times 0 = 3$$

$$\|\mathbf{A}\| = \sqrt{4} = 2, \|\mathbf{B}\| = \sqrt{4} = 2$$

$$\cos(A, B) = 3/(2 \times 2) = 0.75 \checkmark$$

**cos(A, C):**

$$A \cdot C = 1 \times 0 + 1 \times 0 + 1 \times 0 + 1 \times 0 + 0 \times 0 + 0 \times 1 = 0$$

$$\|\mathbf{A}\| = 2, \|\mathbf{C}\| = 1$$

$$\cos(A, C) = 0/(2 \times 1) = 0.00 X$$

**cos(B, C):**

$$B \cdot C = 1 \times 0 + 1 \times 0 + 0 \times 0 + 1 \times 0 + 1 \times 0 + 0 \times 1 = 0$$

$$\|\mathbf{B}\| = 2, \|\mathbf{C}\| = 1$$

$$\cos(B, C) = 0/(2 \times 1) = 0.00 X$$

**Interpretation:** A and B share 75% similarity (both about learning). C is orthogonal (no shared terms).

# Cosine Similarity: Python Implementation

## From Scratch (NumPy)

```
import numpy as np

def cosine_similarity(a, b):
    """Tính cosine similarity giữa hai vectors"""
    dot_product = np.dot(a, b)
    norm_a = np.linalg.norm(a)
    norm_b = np.linalg.norm(b)

    if norm_a == 0 or norm_b == 0:
        return 0.0

    return dot_product / (norm_a * norm_b)

# Ví dụ
D_A = np.array([0.5, 0.3, 0.0, 0.2])
D_B = np.array([0.4, 0.4, 0.0, 0.1])

sim = cosine_similarity(D_A, D_B)
print(f"cos(D_A, D_B) = {sim:.4f}") # Output: 0.9615
```

## Using sklearn

```
from sklearn.metrics.pairwise import cosine_similarity
import numpy as np

# Tạo ma trận documents (mỗi hàng là 1 document)
documents = np.array([
    [0.5, 0.3, 0.0, 0.2], # D_A
    [0.4, 0.4, 0.0, 0.1], # D_B
    [0.0, 0.0, 0.6, 0.3] # D_C
])

# Tính similarity matrix
sim_matrix = cosine_similarity(documents)

print("Similarity Matrix:")
print(sim_matrix)
```

PART 6

# Dimensionality Reduction

---

*"Principal Component Analysis (PCA)"*

From High-Dimensional Sparse → Low-Dimensional Dense

10,000D



300D

# The Curse of Dimensionality

## ⚠ Problem: High-Dimensional Vectors

With  $|V| = 50,000$  words, each vector has **50,000 dimensions!** Most zeros → Sparse and Inefficient

[0, 0, 0.2, 0, 0, 0.1, ...]



Storage  
10K docs × 50K dims = ~2GB



### Computation

All pairs:  $O(n^2d)$  → Very slow



### Sparsity

99% zeros, noise amplification

## ✓ Solution: Dimensionality Reduction

Goals:

- 50,000D → 100-300D
- Preserve semantics

Methods:

- PCA, SVD, LSA

Benefits:

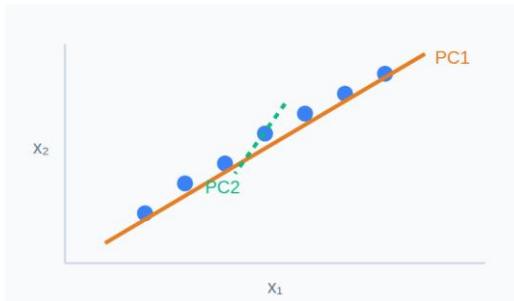
- Faster, less memory

# PCA: Principal Component Analysis

## Definition

PCA finds **orthogonal directions** (principal components) that capture **maximum variance**. Project high-dimensional data onto first k components to reduce dimensions while preserving information.

### Example: 2D → 1D



PC1 captures most variance, PC2 captures remaining

### Key Concepts

- **Principal Components:** Orthogonal directions of max variance
- **Eigenvalues:** Amount of variance captured by each PC
- **Eigenvectors:** Direction of each PC

### Variance Explained

Choose k components that explain ~90-95% of total variance. Scree plot helps visualize: plot eigenvalues and find "elbow".

# PCA: Mathematical Formulation

## Step 1: Center the Data

$$\tilde{X} = X - \mu$$

$\mu_j = (1/n) \sum_i x_{ij}$  (mean of each column)

## Step 2: Covariance Matrix

$$C = (1/(n-1)) \tilde{X}^T \tilde{X}$$

C is  $d \times d$  matrix ( $d$  = original dimensions)

## Step 3: Eigendecomposition

$$Cv = \lambda v$$

v: eigenvector (direction),  $\lambda$ : eigenvalue (variance)

## Step 4: Sort Eigenvectors

Sort eigenvectors by eigenvalues (descending). First eigenvector = direction of maximum variance.

## Step 5: Project to k Dimensions

$$X_{\text{reduced}} = \tilde{X} \cdot V_k$$

$V_k$  = matrix of first  $k$  eigenvectors

## Variance Explained Ratio

$$VE_k = \sum_{i=1}^k \lambda_i / \sum_{j=1}^d \lambda_j$$

Choose  $k$  where  $VE_k \geq 0.90-0.95$

# PCA: Python Implementation

## Using sklearn

```
import numpy as np
from sklearn.decomposition import PCA

# Dữ liệu gốc (4 documents, 1000 features)
X = np.random.rand(4, 1000)

# Khởi tạo PCA với 2 components
pca = PCA(n_components=2)

# Fit và transform
X_2d = pca.fit_transform(X)

print("Shape gốc:", X.shape)      # (4, 1000)
print("Shape sau PCA:", X_2d.shape) # (4, 2)

# Xem variance explained
print("Variance ratio:", pca.explained_variance_ratio_)
print("Total variance:", sum(pca.explained_variance_ratio_))
```

## ⚠ PCA Limitations for Sparse Data

**Problem:** PCA centers data (subtracts mean), destroying sparsity. For TF-IDF matrices, use **TruncatedSVD** instead - it works directly on sparse matrices without centering.

# SVD: Alternative to Eigendecomposition

$$X = U \Sigma V^T$$

**U (n×n)**

Left singular vectors (docs)

**Σ (n×d)**

Singular values (diagonal)

**VT (d×d)**

Right singular vectors (words)

## \_TRUNCATED SVD (LSA)

$$X \approx U_k \Sigma_k V_k^T$$

Keep only k largest singular values. Best rank-k approximation (Eckart-Young theorem).

**LSA:** Apply truncated SVD to TF-IDF matrix → latent semantic space

## sklearn Implementation

```
from sklearn.decomposition import TruncatedSVD
from sklearn.feature_extraction.text import TfidfVectorizer

# Corpus
corpus = [
    "machine learning algorithms",
    "deep learning neural networks",
    "natural language processing",
    "machine learning for NLP"
]

# TF-IDF vectorization
vectorizer = TfidfVectorizer()
X_tfidf = vectorizer.fit_transform(corpus)
print("TF-IDF shape:", X_tfidf.shape) # (4, vocab_size)

# TruncatedSVD
svd = TruncatedSVD(n_components=2, random_state=42)
X_lsa = svd.fit_transform(X_tfidf)

print("LSA shape:", X_lsa.shape) # (4, 2)
print("Variance explained:", svd.explained_variance_ratio_)
print("Total:", sum(svd.explained_variance_ratio_))
```

PART 7

# Applications

---

*"Putting Vector Space Models to Work"*

Document Search

Word Similarity

Document Clustering

# Document Search System

## Document Search Pipeline



## Python Implementation (handout)

# Word Similarity with Vector Spaces

## Finding Similar Words

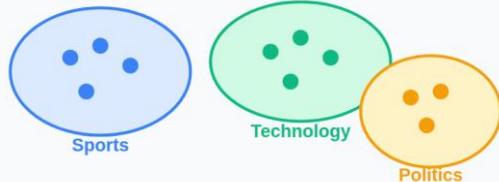
Words in similar contexts have similar vectors. Use cosine similarity to find related words.



# Document Clustering

## Grouping Similar Documents

Use TF-IDF vectors with clustering algorithms to group documents by topic without labels.



### ✓ Pros

Unsupervised, scalable, interpretable clusters

### X Cons

Need to choose k, sensitive to initialization

## Use Cases

News categorization • Customer feedback grouping • Email organization • Research paper clustering

PART 8

# Summary & Lab Practice

---

[Key Concepts](#)

[Formulas](#)

[Lab Practice](#)

[Exercise 1](#)

# Key Concepts Summary

## 1. Word Representations

- One-hot: sparse, no similarity
- Distributional: context-based
- Count vs prediction methods

## 2. Co-occurrence Matrices

- Term-document: document vectors
- Word-word: word vectors
- Window size impacts semantics

## 3. TF-IDF Weighting

- TF: local term importance
- IDF: global term rarity
- Balances frequency vs specificity

## 4. Cosine Similarity

- Measures angle, not magnitude
- Range: [-1, 1] or [0, 1] for TF-IDF
- Cosine distance =  $1 - \text{cosine}$

## 5. Dimensionality Reduction

- PCA: eigendecomposition
- SVD: works on sparse matrices
- LSA: SVD on TF-IDF

## 6. Applications

- Document search: query → rank
- Word similarity: nearest neighbors
- Clustering: unsupervised grouping

# Key Formulas Reference



$$\text{tf-idf}(t,d) = \text{tf}(t,d) \times \log(N/\text{df}_t)$$

$\text{tf}(t,d)$  = count of term t in doc d; N = total docs;  $\text{df}_t$  = docs containing t



$$\cos(A,B) = (A \cdot B) / (\|A\| \times \|B\|)$$

$A \cdot B = \sum a_i b_i$  (dot product);  $\|A\| = \sqrt{(\sum a_i^2)}$  (L2 norm)



$$Cv = \lambda v$$

C = covariance matrix; v = eigen vector;  $\lambda$  = eigen value (variance)



$$X = U\Sigma V^T$$

U = left singular vectors;  $\Sigma$  = singular values; V = right singular vectors



💡 TF-IDF captures term importance • Cosine measures angle • PCA/SVD reduce dimensions

# Lab Practice Overview

## Task 1: TF-IDF Vectors

- Build TF-IDF vectors from scratch
- Compare with sklearn TfidfVectorizer
- Visualize top terms per document

## Task 2: Cosine Similarity

- Implement cosine similarity function
- Build document similarity matrix
- Create similarity heatmap

## Task 3: Document Search

- Build search engine with TF-IDF
- Implement query → TF-IDF → rank
- Evaluate with sample queries

## Challenge

Apply TruncatedSVD to reduce TF-IDF dimensions, then use t-SNE for 2D visualization. Color points by document category to see if similar documents cluster together.

# Exercise 1 Release

## Sentiment Analysis & Vector Spaces

Weight: 10%



End of Session 3



Beginning of Session 5



2 Weeks

### Tasks

#### Part A

Sentiment classifier with Logistic Regression or Naïve Bayes

#### Part B

Build word vectors, compute cosine similarity

#### Part C

Document search with TF-IDF and cosine similarity

### Deliverables

Jupyter notebook with code + explanations • Short report (2-3 pages) with results and analysis